

# Contextual Bandit Portfolio Allocation

Evaluating LinUCB contextual bandit for weekly capital allocation across five technology leaders: AVGO, AAPL, TSLA, GOOGL, NVDA.

## Test Period

52-week evaluation window from September 23, 2024 to September 22, 2025

## Context Features

Lagged returns, volatility, RSI, and relative price/volume signals drive allocation decisions

## Benchmarks

Equal-weight portfolio, seeded random arm selection, SPY ETF, and best single ticker performance



# LinUCB Architecture Overview

*LinUCB feedback loop: feature engineering → state update → UCB policy → realized reward feedback*



## Feature Engineering

Technical indicators and market signals transformed into context vectors for each asset



## State Update

Bayesian parameter estimation updates confidence bounds based on observed rewards



## UCB Selection

Upper confidence bound policy balances exploitation of high-performing arms with exploration



## Reward Feedback

Weekly returns provide reward signal to refine future allocation decisions

# Performance Results

1.91

Ending Value

Portfolio multiplier vs SPY's 2.46x benchmark

90.9%

Total Return

Cumulative performance over 52-week period

2.06

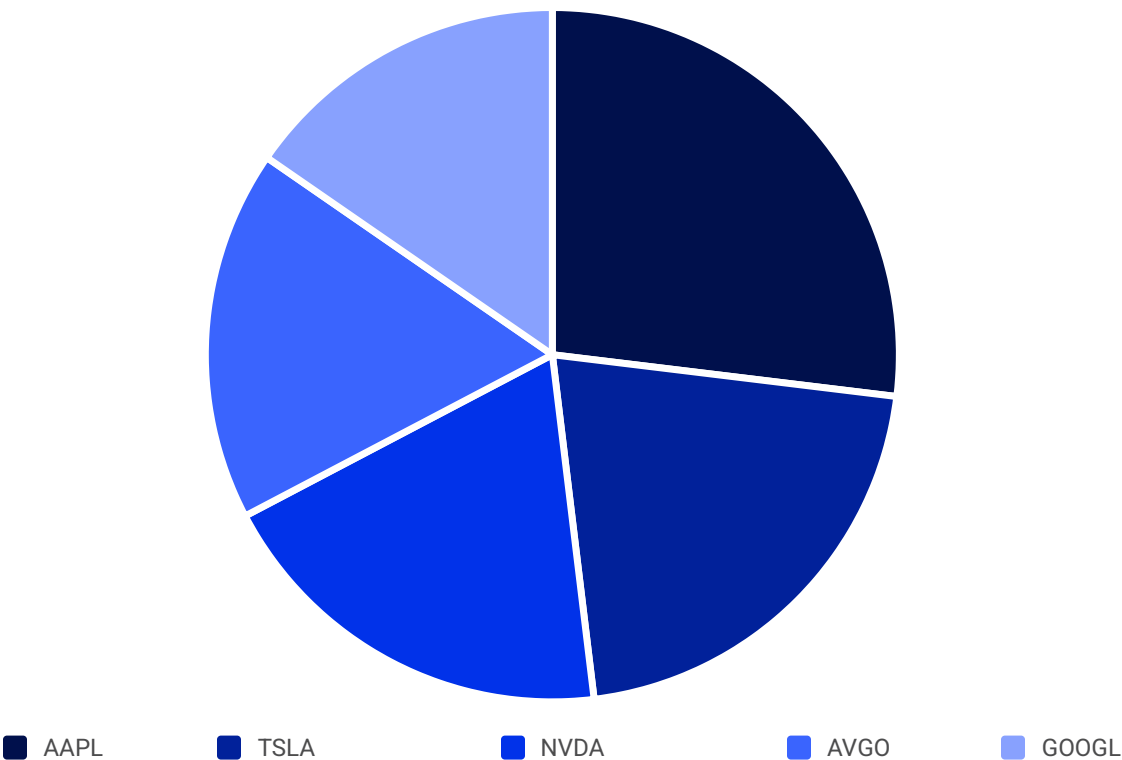
Sharpe Ratio

Risk-adjusted return efficiency metric

-26.7%


Max Drawdown

Peak-to-trough decline during test period




LinUCB demonstrated balanced exploration across all five assets, with AAPL receiving the highest allocation frequency.

# Performance Visualization



## Cumulative Value

Portfolio growth trajectory compared against baseline strategies and SPY benchmark



## Return Resilience

Weekly return distribution showing downside protection and volatility characteristics

Left panel shows cumulative value progression versus baselines | Right panel demonstrates weekly return resilience and risk profile

### Key Insight

LinUCB achieved competitive returns with systematic risk management, demonstrating the viability of contextual bandits for tactical asset allocation in concentrated tech portfolios.

### Implementation Note

The algorithm's balanced exploration-exploitation trade-off prevented over-concentration while capturing momentum signals across the five-asset universe.